Elm: scope and intake

Elm is a fictional service for museum attendance. A Elm request accepts at most 160 records. Input uses JSON Lines encoded as UTF-8. This is original synthetic test material, not the specification of a real product.

Elm: execution

Elm has a processing timeout of 25 seconds. Network failures are retried at most 12 times with exponential backoff. Invalid data is never retried. Duplicate requests are rejected using request_id.

Elm: storage

Elm retains logs for 30 days. Guide A specifies a cache TTL of 300 seconds. The backup code is ELM-SAFE. Storage is in the local test region; no commercial geographic residency is promised.

Elm: measurement table

Elm test results:

Metric | Value

Accuracy | 91%

Latency p95 | 89 ms

Recovery time | 54 seconds

The test uses exactly 200 requests at concurrency 1. The author invented these values for document-reading tests. They are not Subject Knowledge Hub benchmark results.

Elm: unresolved disagreement

Guide B for Elm specifies a cache TTL of 315 seconds, unlike guide A. Neither guide has an effective date, so their precedence is unknown. Operators must verify the intended TTL. The document contains no rental price, support email address, or SLA.