Dataset protocol

The Aurora dataset contains 2400 examples and 12 input features. The split is 80% training, 10% validation, and 10% test. The split random seed is 17. The target has three classes: red, green, and blue.

Training configuration

Aurora uses logistic regression with L2 regularization. The regularization parameter C is 0.5. The maximum number of solver iterations is 200. The solver is lbfgs. No GPU is required for this fixture experiment.

Data cleaning

Duplicate Aurora examples are removed by sample_id before splitting. Missing numeric values use the training median. Missing categorical values use the label unknown. Never fit preprocessing on validation or test examples.

Evaluation protocol

Aurora reports macro F1 as its primary metric because its classes are imbalanced. Macro F1 is 0.81 and balanced accuracy is 0.84. The held-out test split is evaluated once after model selection.

Service limits

The Aurora inference API accepts at most 16 rows per request. Requests time out after 8 seconds. Feature names must match the twelve training columns. The health endpoint is /health. The request schema version is v2.

Reproducibility

The Aurora artifact checksum prefix is a17c9e. The experiment seed is 17 and the project release is 1.2.0. Prediction logs are retained for 7 days. The model must be revalidated whenever the data schema changes.